

SELF-NORMALIZING NEURAL NETWORKS

HURDLES WHEN TRAINING A NEURAL NETWORK

CHRISTOPH LANGE

JANUARY 12, 2021



- 1 Backpropagation
- 2 Issues during the Training Process
- 3 Self-normalizing Neural Networks (SNNs)

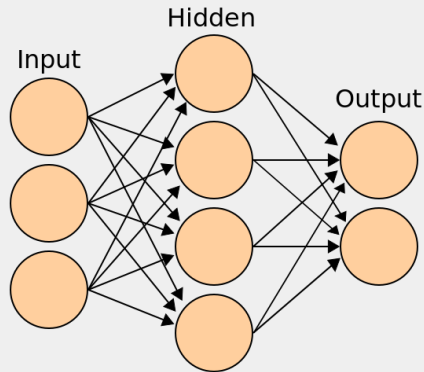
BACKPROPAGATION

TRAINING A NEURAL NETWORK

The paper [Self-Normalizing Neural Networks](#) focuses on Fully Connected Neural Networks. one layer:

- $y(x) = f(Wx + b)$ represents one layer
 - ▶ $x \in \mathbb{R}^n$ activations current layer
 - ▶ $y \in \mathbb{R}^m$ activations next layer
 - ▶ $W \in \mathbb{R}^{m \times n}$
 - ▶ $b \in \mathbb{R}^m$ the bias vector, assume $b = 0$ for simplicity
 - ▶ $f \in C^0$ continous non-linear function
- want to learn W for each layer

Figure: A fully connected neural network (picture from [here](#))



For one layer we have:

$$y(x) = f(z(x)) \quad z(x) = Wx \quad (1)$$

Derivatives with respect to the parameters W and the activations x :

$$\frac{dy}{dW} = \frac{dy}{dz} \cdot \frac{dz}{dW} = f'(z(x)) \cdot \frac{dz}{dW} = f'(Wx) \cdot x^T \quad (2)$$

$$\frac{dy}{dx} = \frac{dy}{dz} \cdot \frac{dz}{dx} = f'(z(x)) \cdot \frac{dz}{dx} = f'(Wx) \cdot W \quad (3)$$

BACKPROPAGATION

For two layers we assume they have the same activation function f :

$$y^{(1)}(x) = f(z^{(1)}(x)) \qquad z^{(1)}(x) = W^{(1)}x \qquad (4)$$

$$y^{(2)}(x) = f(z^{(2)}(x)) \qquad z^{(2)}(x) = W^{(2)}x^{(2)} = W^{(2)}y^{(1)} \qquad (5)$$

The derivatives within one layer are the same:

$$\frac{dy^{(i)}}{dW^{(i)}} = f'(W^{(i)}x^{(i)}) \cdot x^{(i)T} \qquad \frac{dy^{(i)}}{dx^{(i)}} = f'(W^{(i)}x^{(i)}) \cdot W^{(i)} \qquad (6)$$

Across two layers we get:

$$\frac{dy^{(2)}}{dW^{(1)}} = \frac{dy^{(2)}}{dy^{(1)}} \cdot \frac{dy^{(1)}}{dW^{(1)}} = \frac{dy^{(2)}}{dx^{(2)}} \cdot \frac{dy^{(1)}}{dW^{(1)}} \qquad (7)$$

$$\frac{dy^{(2)}}{dx^{(1)}} = \frac{dy^{(2)}}{dy^{(1)}} \cdot \frac{dy^{(1)}}{dx^{(1)}} = \frac{dy^{(2)}}{dx^{(2)}} \cdot \frac{dy^{(1)}}{dx^{(1)}} \qquad (8)$$

Derivative across N layers is :

$$\frac{dy^{(N)}}{dW^{(1)}} = \frac{dy^{(N)}}{dx^{(N)}} \cdot \dots \cdot \frac{dy^{(2)}}{dx^{(2)}} \cdot \frac{dy^{(1)}}{dW^{(1)}} \quad (9)$$

Why is it called Backpropagation?

$$\frac{dy^{(N)}}{dW^{(2)}} = \underbrace{\frac{dy^{(N)}}{dx^{(N)}} \cdot \dots \cdot \frac{dy^{(3)}}{dx^{(3)}}}_{=: u^{(2)}} \cdot \frac{dy^{(2)}}{dW^{(2)}} = u^{(2)} \cdot \frac{dy^{(2)}}{dW^{(2)}} \quad (10)$$

From the perspective of layer 2, $u^{(2)}$ is the derivative that comes from up front. And the components are travelling backwards.

ISSUES DURING THE TRAINING PROCESS

PROBLEM SETUP

This is the [Fashion MNIST dataset](#)

- contains of 70000 cloth items
- one picture and one label per item, i.e. "Dress", "Sneaker", "Coat"
- each picture is gray-scaled and 28 x 28 pixels
- example use case for now
- similar to classifying bacteria
- complex problem needs a more complex deeper network



EXAMPLE

Let's oversimplify:

- each layer consists of one neuron
- the derivative of each layer is $\frac{dy^{(i)}}{dx^{(i)}} = c$ constant
- we have $N = 10$ layers

What happens to the derivative of the first layer parameters $W^{(1)}$?

$$\begin{aligned}\frac{dy^{(10)}}{dW^{(1)}} &= \frac{dy^{(10)}}{dx^{(10)}} \cdot \dots \cdot \frac{dy^{(2)}}{dx^{(2)}} \cdot \frac{dy^{(1)}}{dW^{(1)}} \\ &= c \cdot \dots \cdot c \cdot \frac{dy^{(1)}}{dW^{(1)}}\end{aligned}$$

VANISHING / EXPLODING GRADIENT PROBLEM

$$\frac{dy^{(N)}}{dW^{(1)}} = c^{N-1} \cdot \frac{dy^{(1)}}{dW^{(1)}}$$

For a huge number of layers N :

- $\frac{dy^{(N)}}{dW^{(1)}}$ get really big for $c > 1$
- for $c \approx 1$ the derivative $\frac{dy^{(N)}}{dW^{(1)}} \approx \frac{dy^{(1)}}{dW^{(1)}}$
- when $c < 1$ the gradient vanishes $\frac{dy^{(N)}}{dW^{(1)}} \approx 0$

For more general cases of multiplying the Jacobians

$$\frac{dy^{(N)}}{dW^{(1)}} = \frac{dy^{(N)}}{dx^{(N)}} \cdot \dots \cdot \frac{dy^{(2)}}{dx^{(2)}} \cdot \frac{dy^{(1)}}{dW^{(1)}}$$

the largest eigenvalues determine the convergence behavior.

SELF-NORMALIZING NEURAL NETWORKS (SNNs)

Let us look at one layer:

- $y(x) = f(Wx)$
- $x \in \mathbb{R}^n$ activations current layer
- $y \in \mathbb{R}^m$ activations next layer
- $W \in \mathbb{R}^{m \times n}$
- $f \in C^0$ continuous non-linear function
- mean of the activations $\mu = \mathbb{E}(x_i)$
- variance of the activations
 $\nu = \text{Var}(x_i)$
- let w^i be the i -th row of W
- the mean of the i -th row weights is
 $\omega^i := \sum_j w_{ij}$
- accordingly the second moment
 $\tau^i := \sum_j w_{ij}^2$

How do mean μ and variance ν of the activations change to the next layer?

$$\begin{pmatrix} \mu \\ \nu \end{pmatrix} \mapsto \begin{pmatrix} \tilde{\mu} \\ \tilde{\nu} \end{pmatrix} = g \begin{pmatrix} \mu \\ \nu \end{pmatrix}$$

Definition

A neural network is called **self-normalizing** if there is a domain

$\Omega = \{(\mu, \nu) \in \mathbb{R} \mid \mu \in [\mu_{\min}, \mu_{\max}], \nu \in [\nu_{\min}, \nu_{\max}]\}$ and a mapping $g : \Omega \rightarrow \Omega$ such that

- $g(\Omega) \subset \Omega$ is a contraction
- g has a stable and attracting fixpoint $(\mu^*, \nu^*) \in \Omega$

SELF-NORMALIZING MAP

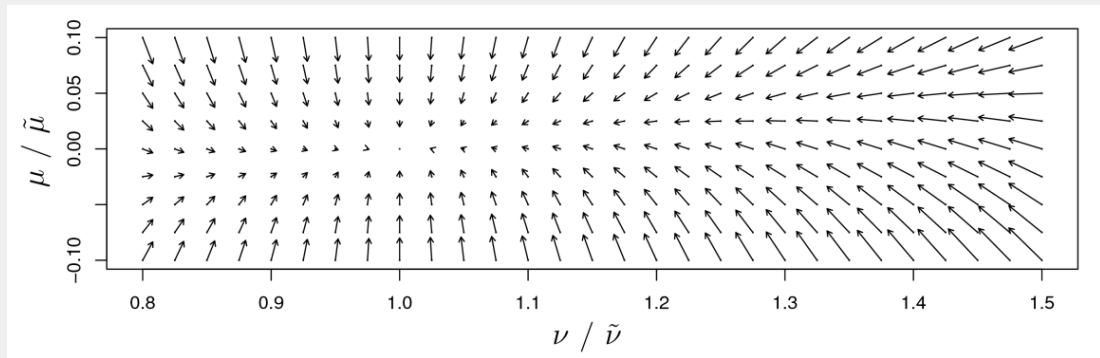


Figure: Assuming the weight of each row w^i to be normalized ($\omega = 0, \tau = 1$), we observe an attracting fixpoint ($\mu^* = 0, \nu^* = 1$). We need to pick the activation function accordingly.

CONSTRUCTING A SELF-NORMALIZING NEURAL NETWORK

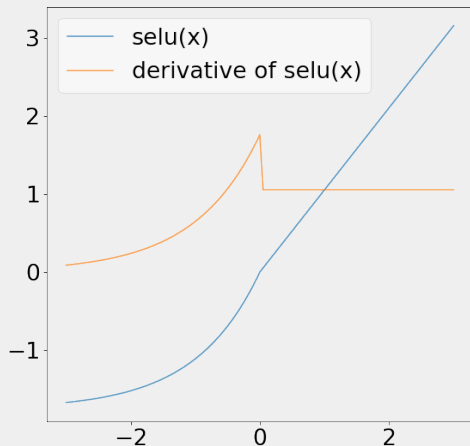
The key component to a SNN is to pick the activation function $f(x)$ to be

$$\text{selu}(x) = \lambda \begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases} \quad (11)$$

a **scaled exponential linear unit**

Assuming

- normalized weights: $\omega = 0, \tau = 1$
- $z(x)$ to be normally distributed specifies $\alpha \approx 1.6733$ and $\lambda \approx 1.0507$.

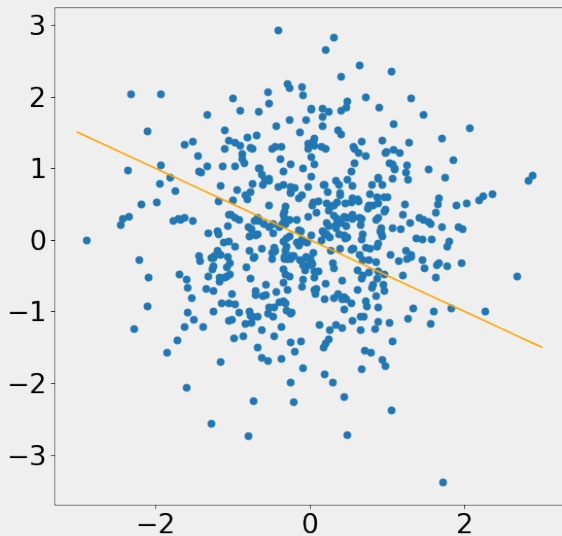


WHY DOES THAT HELP?

- supports deeper fully connected neural networks
- matches a weight initialization of $\omega = 0$ and $\tau = 1$ very well
- automatically normalizes the activations of each layer
 - ▶ $y^{(n)} \circ y^{(n-1)} \circ \dots \circ y^{(1)}(x^{(1)})$ with normalized input features $x^{(1)}$
 - ▶ $y^{(n)} \circ y^{(n-1)} \circ \dots \circ y^{(k)}(x^{(k)})$ with normalized activations $x^{(k)}$
- prevents the gradient from exploding
 - ▶ derivative of $\text{selu}(x)$ is highly bounded
 - ▶ extreme activations are damped by the convergence to $(\mu, \nu) = (0, 1)$
- guarantees non vanishing gradients

THANK YOU
QUESTIONS?

[LINK TO SLIDES](#)



REFERENCES



GÜNTER KLAMBAUER, THOMAS UNTERTHINER, ANDREAS MAYR, AND SEPP HOCHREITER.
SELF-NORMALIZING NEURAL NETWORKS, 2017.